

UW Tower Data Center

Research Computing Club Presents

# Hyak Training Session

January 24, 2022

Brenden Pelkie (Training Coordinator)

# Outline

Where I assume you're at right now:

"I finished the steps to getting access\* to Hyak, but don't really know what to do now. How do I get up and running with the computing I want to do?"

## I. Hyak overview

- A. Hyak architecture
- B. Logging on to Hyak

## II. Navigating Hyak

- A. Basic commands
- B. Important locations
- C. Transferring files

## III. Job Scheduling - Slurm

- A. Running a job
- B. Commands
- C. Modules
- D. Interactivity

## IV. Other resources

- A. Topics outside of this tutorial
- B. Places to get help

\* <https://depts.washington.edu/uwrcc/getting-started-2/getting-started/>

# Things We Won't Cover:

By the end of this training, you will not:

- Know how to build and install software on Hyak
- Have a complete workflow for your particular project
- Understand how to optimize your code for an HPC system
- (necessarily) know/have the best tools for your problem:  
There are lots of ways of doing most things I will present here.

# What is high performance computing?

Real research problems require large calculations to solve

Compared to your personal computer, HPC has:

- More processor cores (40 vs 4-8)
- More memory
- Multiple nodes (ie multiple computers)
- Faster networking for better parallelization

Examples: Summit, Fugaku, Perlmutter, etc.

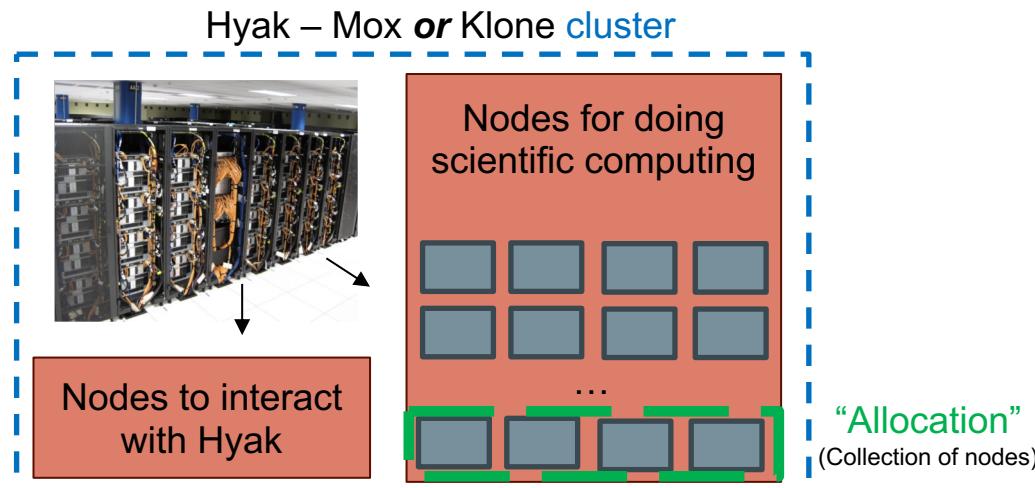


# What is Hyak?

Hyak is UW's research computing environment

Comprised of two supercomputing "clusters":

- Mox - in phase-out process
- Klone – current cluster



Nodes on Mox and on Klone are only accessible if signed onto that [cluster](#)

On Hyak, jobs are run on *nodes*:

- physically an individual computer/server
- Can use part of a node or multiple nodes for a job

## Mox nodes:

- 28 cores
- 128 GB RAM

## Klone nodes:

- 40 cores (can be split between users)
- 192 GB RAM (or 768 GB)

# Resource Allocations on Mox and Klone

Access to nodes is through a specific allocation or account

**stf:** All tuition-paying students have access to the student technology fee ('stf') allocation

- Ideally for use on personal or class work – **not for research projects**
- Mox: 54 compute nodes, 9 GPU nodes
- Klone: 15 compute nodes, 1 GPU node, 3 hugemem nodes

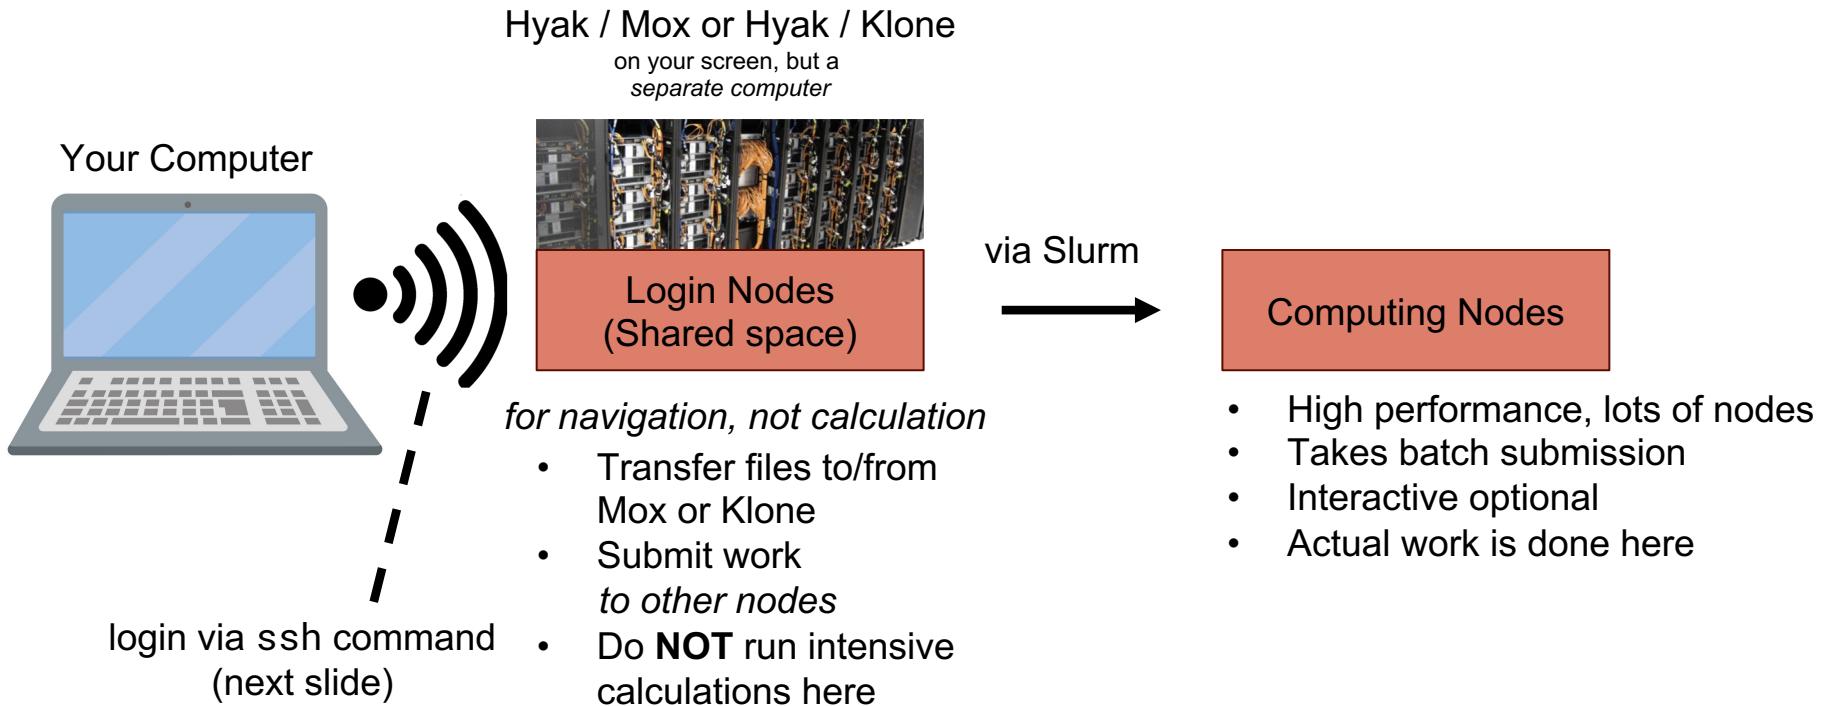
**Research group allocation:** Purchased by advisor. Use this for research.

**checkpoint:** allows you to use idle resources from other allocations. Job may be terminated at any time – plan accordingly

Use `hyakalloc` to view available resources

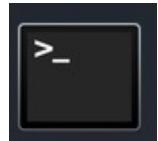
# Accessing Nodes - Hyak “Architecture”

- All nodes on same cluster share the same filesystem



# Logging on to Hyak

## 1. Launch a new terminal



Terminal – “command line”

Mac/Linux: you have this

Windows: Hang tight

```
home -- bash -- 43x7
Last login: Mon Oct 26 13:03:34 on ttys004
(base) Jesses-MacBook-Pro:~ home$
```

## 2. Log into Hyak with SSH

“secure shell” command

Enables connection to remote machine

```
(base) Jesses-MacBook-Pro:~ home$ ssh -X prelj@mox.hyak.uw.edu
Password: ?
```

ssh <uwnetid>@**mox**.hyak.uw.edu

ssh <uwnetid>@**klone**.hyak.uw.edu

## 3. Enter your UW Credentials



Two-factor authentication required  
(Through UW-IT)

## 4. You're in!

```

This login node is meant for interacting with the job scheduler and
transferring data to and from Hyak. Please work by requesting an
interactive session on (or submitting batch jobs to) compute nodes.

Visit the Hyak user wiki for more details:
http://wiki.hyak.uw.edu/Hyak+mox+Overview

Questions? E-mail help@uw.edu with "hyak" in the subject.

Run "scontrol show res" to see any reservations in place that will
prevent your jobs from running with a "(ReqNodeNotAvail,*)" error.

(base) [prelj@mox1 ~]$
```

- Notice you are defaulted to the *login node*: **Jobs are not run here**

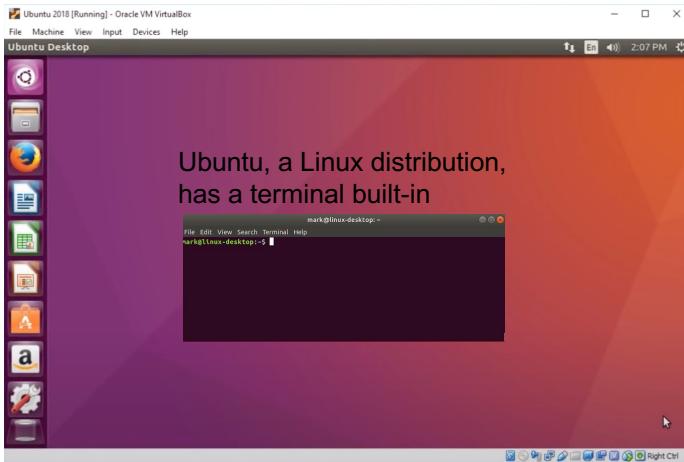
# What if I'm on Windows?

You'll need a terminal somehow, and Windows does not provide a good one by default  
RCC recommendation: Windows Subsystem for Linux\* (I am not a windows user)

- Allows you to run a linux shell in Windows
- Access windows filesystem from linux shell (useful for transferring data to/from hyak)

To install on a recent Windows version:

1. Launch Powershell and enter command “wsl –install”
2. After install completes, reboot
3. Open linux distribution from start menu and set up profile
4. Use SSH to access Hyak, as described on last slide



Other options for Windows:

- PuTTY (Terminal emulator)
- Linux virtual machine (Similar concept to WSL)

More help:

<https://docs.microsoft.com/en-us/windows/wsl/install>

<http://wiki.cac.washington.edu/display/hyakusers/Logging+In>

# Practice: Log In

1. Open a new terminal window
2. Enter the command “ssh <your netid>@klone.hyak.uw.edu
3. Follow the prompts for your password (same as UW password) and 2FA
  - Note: As you type your password, nothing will appear. This is normal.



**I'm in**

... Now what?

# Exploring Hyak/Command line 101



Hyak works on a command line interface: no mouse and pointer

- Commands are written in *bash*
- Basic commands transferable to all linux/mac systems

```
(base) [prelj@mox1 ~]$
```

## Essential bash commands for navigation:

Note: On a computer, a “folder” and a “directory” are the same thing

### Where in the computer am I?

`pwd` - “Print working (current) directory”

```
(base) [prelj@mox1 ~]$ pwd  
/usr/users/prelj
```

My home directory (see next slide)

### Change to a different directory:

`cd <filepath>` - “Change directory”

```
(base) [prelj@mox1 ~]$ cd /gscratch  
(base) [prelj@mox1 gscratch]$
```

`cd ..` - “Move up one directory”

### What folders and files are here?

`ls` - “List” files in current directory (folder)

```
(base) [prelj@mox1 ~]$ ls  
d g_perf intel opt
```

### What does this file say?

`cat <filename>`

```
(base) bgpelkie@klone1:~$ cat .bash_profile
```

# More shell (bash) commands: Try them out!

## File system manipulation:

- `Mkdir <dir. name>`
  - “Make directory”
- `Mv <filename> <new location>`
  - Move (rename) file or directory
- `Cp <filename> <new location>`
  - “Copy” files and/or directories (-r)
- `rm <filename>`
  - “Remove” files and/or directories (-r)
  - Add “-r” between rm and filename for directories
  - This deletes the file permanently. Be careful. Have a backup of important stuff.



There are many more – here's one resource:  
<https://guide.bash.academy/>

## File editing:

Edit text files (ie code, scripts) with a command line text editor

Pick one and get comfy with it

- Vim
- Nano
- emacs

## Changing Bash Profile

You can change behavior of commands (ie get more information from `ls`) and add colors by editing `.bash_profile` file in home directory. See files from this training for a starting point

# Important locations on Hyak (get there with cd)

## ★ /gscratch/stf

- Main work location for stf users
- **Any files untouched for >30 days will be auto-deleted!!**

## ★ ~

- Home directory “~”, *you are put here by default*
- /usr/lusers/<username> on Mox
- /mmfs1/home/<username> on Klone
- **Only 10 GB of storage per user**

## ● /tmp

- Node local storage (separate from shared filesystem)

## ● /sw

- All software installs

## ● /sw/contrib

- User installed software – *not yet operational on Klone*

## ● /sw/modules-1.775/modulefiles/contrib

- User added modulefiles – on mox, still not up for Klone



Where pre-installed  
programs (and their  
supporting files) are housed

You will seldom  
need to go here

# Transferring files

scp – Secure Copy

- Send file To Mox or Klone:

scp <path/to/file> <username>@**mox**.hyak.uw.edu:<path/to/dest>



On your computer

...@**klone**.hyak...



Your login from ssh



Place on Mox you want  
to put the file

- Get file From Mox:

scp <username>@**mox**.hyak.uw.edu:<path/to/file> <path/to/dest>  
...@**klone**.hyak...

Flag **-r** needed to copy whole folder at once

An alternative copier for synchronizing directories is **scopy**

# Running Your First Job

Real computing jobs are passed to Slurm

- Ensures fair share between users
- Manages multi-node jobs

Interactive vs. Batch jobs:

- Interactive: requires user input during job
- Batch: job runs unattended

Some information is needed during job setup:

Partition: whose CPUs/GPUs are we using? (we'll use stf)

Allocation: which bank account pays for these? (also stf)

Locations: where to write/read files to/from

Resources: how many CPUs/GPUs are needed, and for how long?

Modules: which programs should be loaded to use?



Packaged into a  
“Slurm script”

# Batch Jobs use a “slurm script”

In a text file called “submit.slurm” on **Klone**

```
#!/bin/bash
## The first line has to say this, as a bash script

## Job Name
#SBATCH --job-name=test_python

## Partition and Allocation
#SBATCH -p compute
#SBATCH -A stf

## Resources
#SBATCH --nodes=1
#SBATCH --time=0:01:00
#SBATCH --ntasks=1
#SBATCH --mem=100G

## Specify the working directory for this job
#SBATCH --chdir=.

## Import any modules here
module load contrib/anaconda/anaconda4.4.0

## Scripts to be executed here
python test_run.py

## Clean up
exit 0
```

Any line that begins with ## is a comment and is not read  
Each line that begins with #SLURM specifies info for slurm  
The #SLURM options can be provided in any order

- Name is optional
- Allocation: Use stf funding/budget
- Partition: *Different on Klone!* “compute” or “ckpt”
  
- Request (1) node
- Time – Hours : Minutes : Seconds
- **Be mindful of your node capacity:** 28 vs. 40 core, memory
  
- Directory “.” is *current location* to read/write files (the same place submit.slurm is saved)
- Load in a pre-installed python package
  
- The real run command

To execute this job,

**sbatch submit.slurm**

# Then We Wait

- `sbatch <script>`
  - Submits a script for non-interactive use
  - Used this to submit a job on previous slide

Want to check the status?

- `sacct`
- `squeue`
  - Flag `-u <uwnetid>` for only your jobs
  - Specify `-p` or `-A` for whole queue

What it looks like to check queue:

(base) [prelj@mox1 move]\$ squeue -p stf	JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
	417044	stf	test_pyt	prelj	CG	0:02	1	n2282
	416180	stf			PD	0:00	1	(Dependency)
	416186	stf	Other peoples' job names and IDs		PD	0:00	1	(Dependency)
	413137	stf			R	1-04:27:53	10	n[2143, 2151, 2166, 2168, 2171-2172,
	415317	stf			R	15:51:35	6	n[2140, 2165, 2275, 2277, 2292, 2298]
	414068	stf			R	22:51:51	6	n[2266, 2269-2270, 2285, 2290, 2308]
	413181	stf			R	1-04:04:16	6	n[2139, 2144, 2146, 2279, 2284, 2312]

Something wrong? / Taking too long?

- `scancel <jobid>`
  - Cancels an unfinished job
  - Can only cancel your own

*Always check your output!*

JOBID is used to refer to specific jobs

NODELIST has ID #s for individual nodes being used

ST column is status

R – Running (currently active)

PD – Pending (in queue)

CG – Completing (in process of termination)

# Practice: Submit a job

While logged into kclone:

1. Change directories to gscratch: “`cd /gscratch/stf`”
2. Make a directory for yourself in gscratch: “`mkdir <your netid>`” and switch into that directory
3. From a new terminal on your computer, copy the handout files to your new gscratch directory: “`scp <path to handout files>/* <your netid>@kclone.uw.edu:/gscratch/stf/<your netid>`”
4. Submit the FizzBuzz job: “`sbatch slurmscript.sh`”
5. View the output file: “`cat <output file name>`”

# Check Memory Resources: Hyakalloc

On Mox:

Allocation Name	Nodes	Cores	Total Memory	Max Mem Per Node
alys	25	40	500G	500G
anantram		40	374G	G
ark		40	373G	G
ark-gpu	1	40	373G	373G
astro	12	336	4992G	500G
baker	25	700	6136G	248G
barnardlab	24	40	185G	248G
bdata		40	373G	G
bdata-gpu	1	40	373G	373G
bkrs	65	1820	16120G	248G
brunton	25	40	752G	752G
build	2	56	496G	248G
build-gpu	1	40	373G	370G
build-long			G	G
bumblerem	7	196	980G	248G
cardss	4	128	740G	185G

... (many!)

On Kfone:

Account-Partition - Resource	Total	Used	Free
chem			
- CPU cores	480	0	480
- Memory (GB)	2220	0	2220
Account-Partition - Resource	Total	Used	Free
pfaendtner			
- CPU cores	400	288	112
- Memory (GB)	1850	260	1590
Account-Partition - Resource	Total	Used	Free
stf			
- CPU cores	680	59	621
- Memory (GB)	3145	900	2245
Account-Partition - Resource	Total	Used	Free
stf-hugemem			
- CPU cores	120	80	40
- Memory (GB)	2256	1500	756

Nodes can be automatically shared between users on Kfone

Only cores are reported

# Running Interactive Jobs from Command Line

```
srun -p compute -A stf --ntasks=8 --mem=10G --time=0:10:00 --pty bash -l
```

Time limit in place for interactive nodes: 8 hrs

- Partition
  - Account
  - Number of processes (\*)
  - Amount of RAM
  - Time
  - Command
- 
- Much of the same specifications as submit.slurm
  - `--pty bash -l` signifies *interactive*
  - Now you have command line control *while* accessing a compute node

Before running interactive job command

```
[(base) [prelj@mox1 jesse]$
```

After running interactive job command

```
[(base) [prelj@n2232 jesse]$
```

- Interactive nodes are computing nodes. You *can* submit jobs to other nodes from interactive nodes.

<https://slurm.schedmd.com/srun.html>  
man srun

# Loading software: the modules system

- **module avail**
  - Show all available modules
- **module load <module>**
- **module unload <module>**
  - (Un)load a given module
  - Provide full module name
- **module list**
  - List loaded modules
- **module purge**
  - Unload all modules
- **module help**
  - Print help with commands

## Running module avail

```
(base) [prel]@mox1 jesse$ module avail
                               /sw/modules-1.775/modulefiles -----
anaconda2_4.3.1                                         contrib/multiz/20090121
anaconda2_5.3                                         contrib/multiz/20190527
anaconda3_4.3.1                                         contrib/mummer/3.23
anaconda3_5.3                                         contrib/mummer/4.0.0b2
ansys_18                                              contrib/muscle/3.8.31
cmake/3.11.2                                         contrib/mysql/8.0.11
cmake_3.8                                            contrib/NAMD2019/namd2
contrib/3ddna/170123                                     contrib/naughty26r11/naughty26r11
contrib/9.0                                             contrib/ncbi-blast/2.7.1
contrib/abaqus/2020                                      contrib/ncbi-vdb/2.9.0
contrib/abyss/2.0.3                                       contrib/nekstar
contrib/adapter-removal/2.1.7                           contrib/newton-x
contrib/admixtools/1.0.1                                contrib/newtonX
contrib/admixture/1.3.0                                 contrib/ngs/2.9.0
contrib/anaconda-2019.03                               contrib/ngsRelate/1.0
contrib/anaconda/2-5.0.1                                contrib/ngsRelate/2.0
contrib/anaconda/3-4.4.0                                contrib/nvidia_toolkit/9.0
contrib/anaconda/3-5.3.0                                contrib/nvidia_toolkit/10.0
contrib/anaconda/anaconda4.4.0                          contrib/NWChem/NWChem
```

... and more

- 400 modules available on Hyak Mox right now!
  - Far fewer on Klone at the moment – hang tight
- Useful programs you'd use on your own computer, i.e.
- Python/Apache, Gromacs, NAMD, R, Mathematica, Matlab, Gaussian, compilers and more
- Parallel computing tools
- CUDA, IMPI, etc.

## Advanced user's note:

The modules system works by keeping track of and modifying environment variables (e.g. PATH, LD\_LIBRARY\_PATH, CPATH, etc.)  
<https://modules.readthedocs.io/en/latest/>

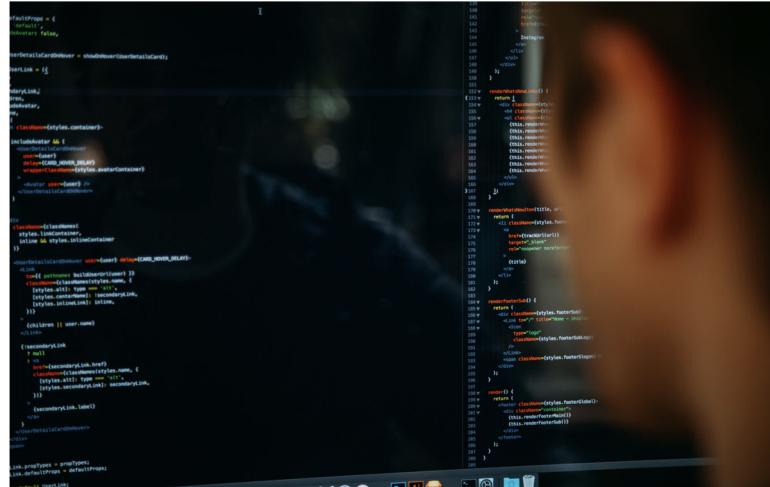
# Checkpoint queue – special case

The checkpoint queue allows any user to run on other groups' unused nodes!

- Partition: ckpt
- Account: <group>-ckpt (stf-ckpt)
- Jobs can be interrupted at any time
- Jobs *will* be interrupted after 4 contiguous hours
- Jobs will be resubmitted after interruption (if under total requested time limit)
- **Your code must be checkpointerd to take advantage of this**
  - Save a binary file containing state of some objects, restart from logs, etc.
  - Your script should also account for any checkpointing that is done

# Topics not covered

- How to install new software
  - Installation and build systems
  - Writing your own modulefiles
- How to parallelize your code
- Anything with the cloud
- General architecture of HPC systems
  - Interconnectivity and node locality
  - Physical infrastructure
- Archiving on Lolo



```
class FileProcessor(object):
    def __init__(self, path, mode='r', encoding='utf-8'):
        self.path = path
        self.mode = mode
        self.encoding = encoding
        self._content = None
        self._is_file = os.path.isfile(path)

    @property
    def content(self):
        if self._content is None:
            with open(self.path, self.mode, encoding=self.encoding) as f:
                self._content = f.read()
        return self._content

    @content.setter
    def content(self, value):
        if self._is_file:
            with open(self.path, self.mode, encoding=self.encoding) as f:
                f.write(value)
        else:
            raise ValueError("Cannot write to a directory")

    def read(self):
        return self.content

    def write(self, content):
        self.content = content

    def list(self):
        if self._is_file:
            raise ValueError("Cannot list files in a file")
        else:
            return os.listdir(self.path)

    def __enter__(self):
        return self

    def __exit__(self, exc_type, exc_value, traceback):
        if self._is_file:
            self.close()

    def close(self):
        pass

    def __str__(self):
        return f"FileProcessor({self.path})"

    def __repr__(self):
        return str(self)

    def __eq__(self, other):
        if isinstance(other, FileProcessor):
            return self.path == other.path
        return False

    def __ne__(self, other):
        if isinstance(other, FileProcessor):
            return self.path != other.path
        return True

    def __hash__(self):
        return hash(self.path)

    def __len__(self):
        if self._is_file:
            return len(self.content)
        else:
            return len(os.listdir(self.path))

    def __iter__(self):
        if self._is_file:
            return iter(self.content)
        else:
            return iter(os.listdir(self.path))

    def __next__(self):
        if self._is_file:
            return next(self.content)
        else:
            return next(os.listdir(self.path))

    def __bool__(self):
        return self._is_file
```

# Where to get help

- Documentation (man or webpages)
- Hyak wiki:  
<https://wiki.cac.washington.edu/display/hyakusers/WIKI+for+Hyak+users>
- Slack channel: <https://uw-rcc.slack.com/>
- Website: <https://depts.washington.edu/uwrcc/>
- Emails: [hpcc@uw.edu](mailto:hpcc@uw.edu) or [uwrcc@uw.edu](mailto:uwrcc@uw.edu)
- Office hours: Mondays 1-2 PM this quarter, location varies

# Happy computing!



# Slurm: Commands

- `sbatch <script>`
  - Submits a script for non-interactive use
- `squeue`
  - Get status of jobs in batch queue
- `scancel <jobid>`
  - Cancels an unfinished job
- `srun`
  - Submits job for interactive use or initiate job steps inside batch script
  - (More on interactive jobs in a minute)
- `sinfo`
  - Get state of partitions and nodes (is the system operational)
- `sacct`
  - Gets accounting information about active and completed jobs

I use these  
most often

**Slurm docs and `man <slurm-command>` are very useful!**  
("man" for manual)